




# Indexing

# Overview

- Indices
  - Tree indices
    - B-Tree
    - Bitmap
  - Composite
- Administration
- Strategy
- Query Explain

# Indexing



# What is an Index?

- An *index* is a optional data structure that allows us to map a key value to a physical location
  - ❑ Requires additional storage
  - ❑ Requires additional accesses for retrieval
  - ❑ Must be updated with database
- Provides “indirect” random access via an index table/structure to any tuple
- Performance is enhanced if index fits into main memory

# Types of Indices

- A *primary index* is based on key (unique) attributes that are used to physically order the data.
- One or more *secondary indexes* can be based on nonordering attributes, both unique and nonunique.
- B-Tree (default)
- Bitmap - store rowids associated with a key value as a bitmap
- Sparse Index
  - ❑ Key and pointer for every *block*
  - ❑ Pointer is usually *direct* to record
- Dense Index
  - ❑ Key and pointer for every *record*
  - ❑ Pointer is usually *indirect* (uses primary index)

# Cluster

- Method for storing more than 1 table per block
- i.e.  
    SELECT last\_name, department\_name  
        FROM Employee JOIN Department USING dept\_id;
- Cluster this data by dept\_id
  - ❑ All the rows in Employee for dept\_id = 10 will be stored with all the rows from Department = 10
  - ❑ 1 I/O to get all data
  - ❑ Create cluster key – dept\_id
  - ❑ Index on cluster key becomes Cluster Index

# Clustered vs Non-Clustered

- Clustered index
  - ❑ Physically reorders the way the records are stored in the table
  - ❑ One per table
  - ❑ Leaf nodes contain the data pages
  - ❑ PK creates if no other clustered index exists
- Non-clustered
  - ❑ Logical order of rows does not match physical order of rows
  - ❑ Leaf nodes contain index rows

---

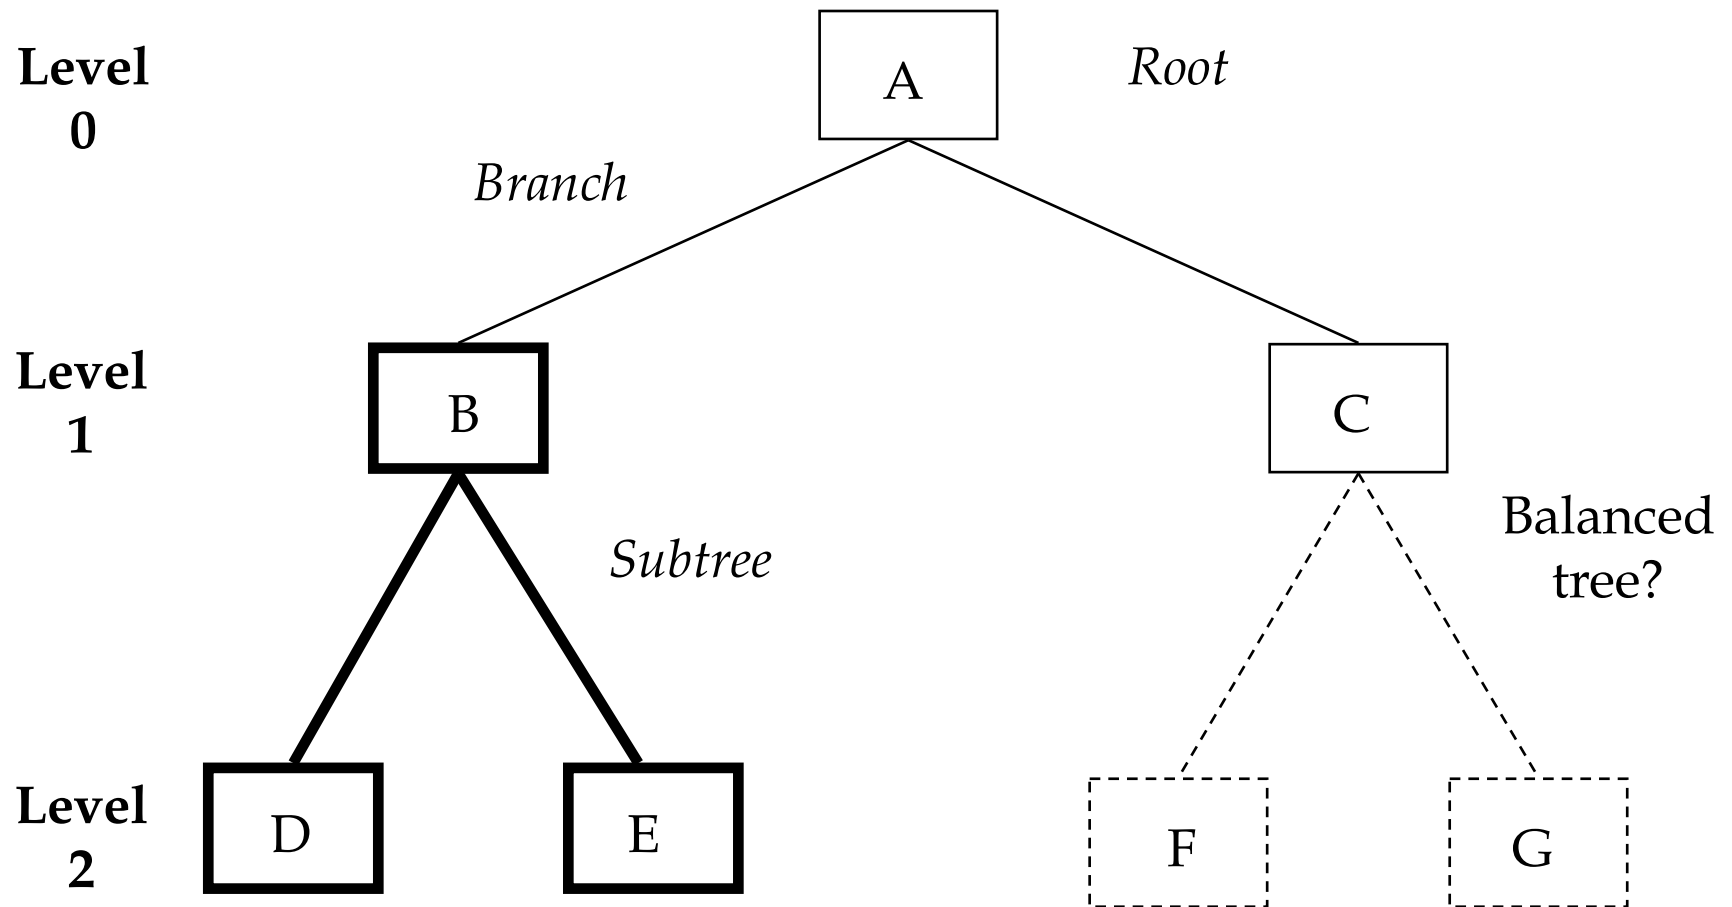
# Tree Structured Indices



# Tree Structured Indices

- *Rooted tree* – forms a hierarchy of index records
- *Leaf nodes* – lowest-level (terminal) nodes
- *Siblings* – nodes that share a common *parent*
- *Path* – set of pointers (branches) from one node to another
  - 1 unique path to each node
- *Degree* – number of siblings permitted
- *Depth* (height) – number of levels
  - root – level zero

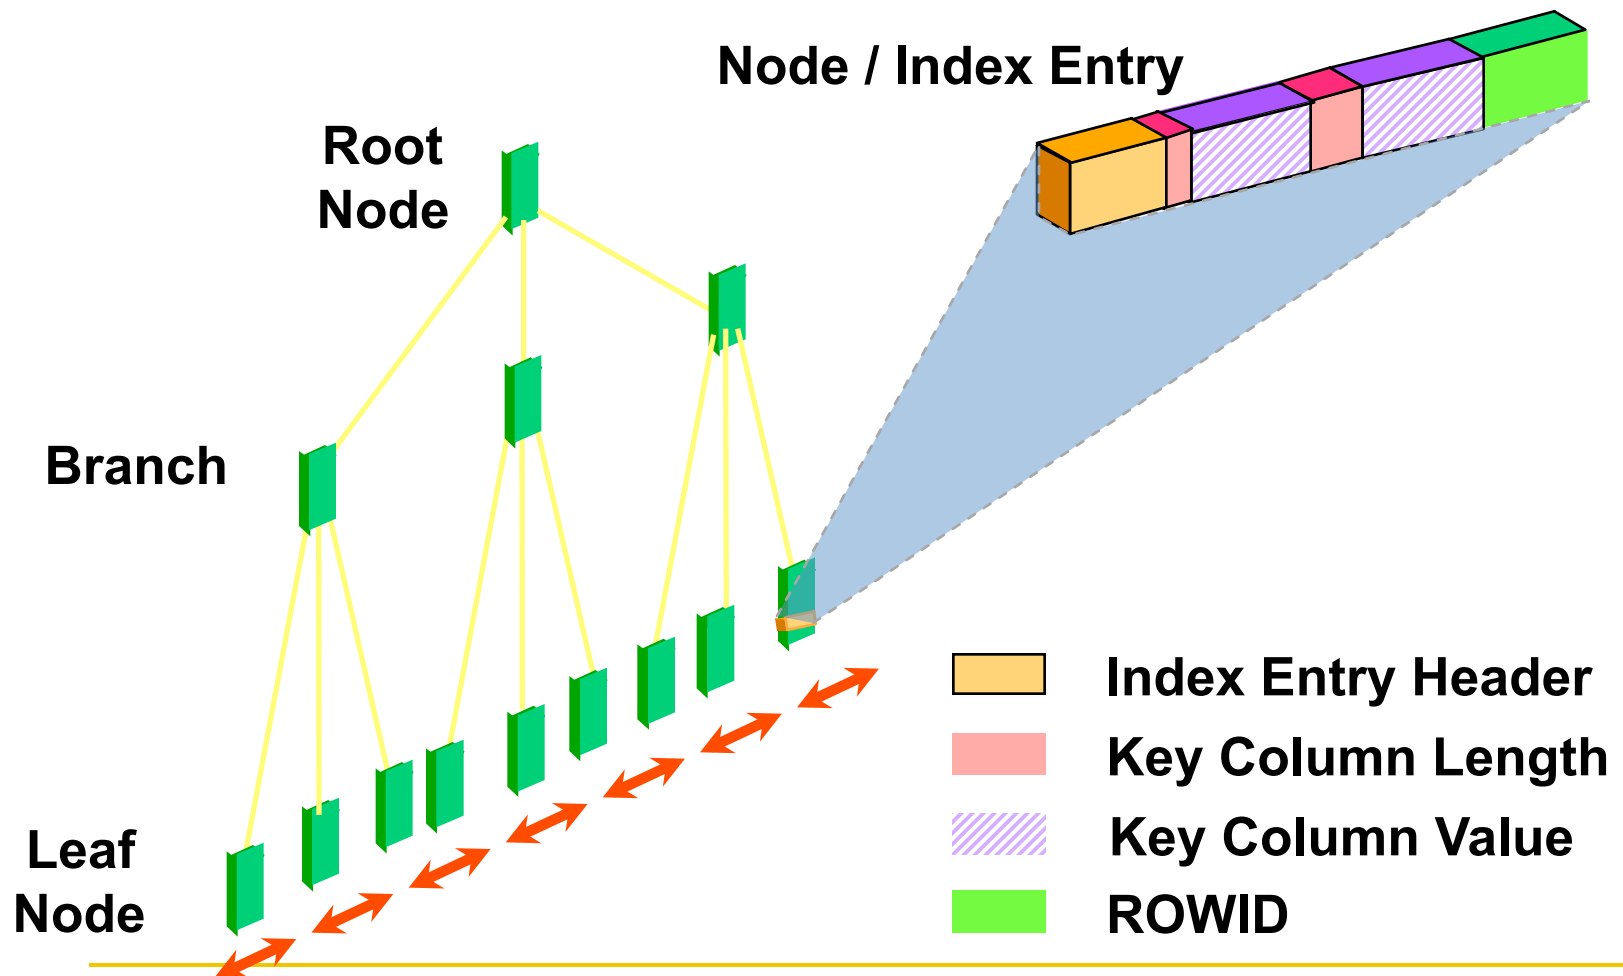
# Trees



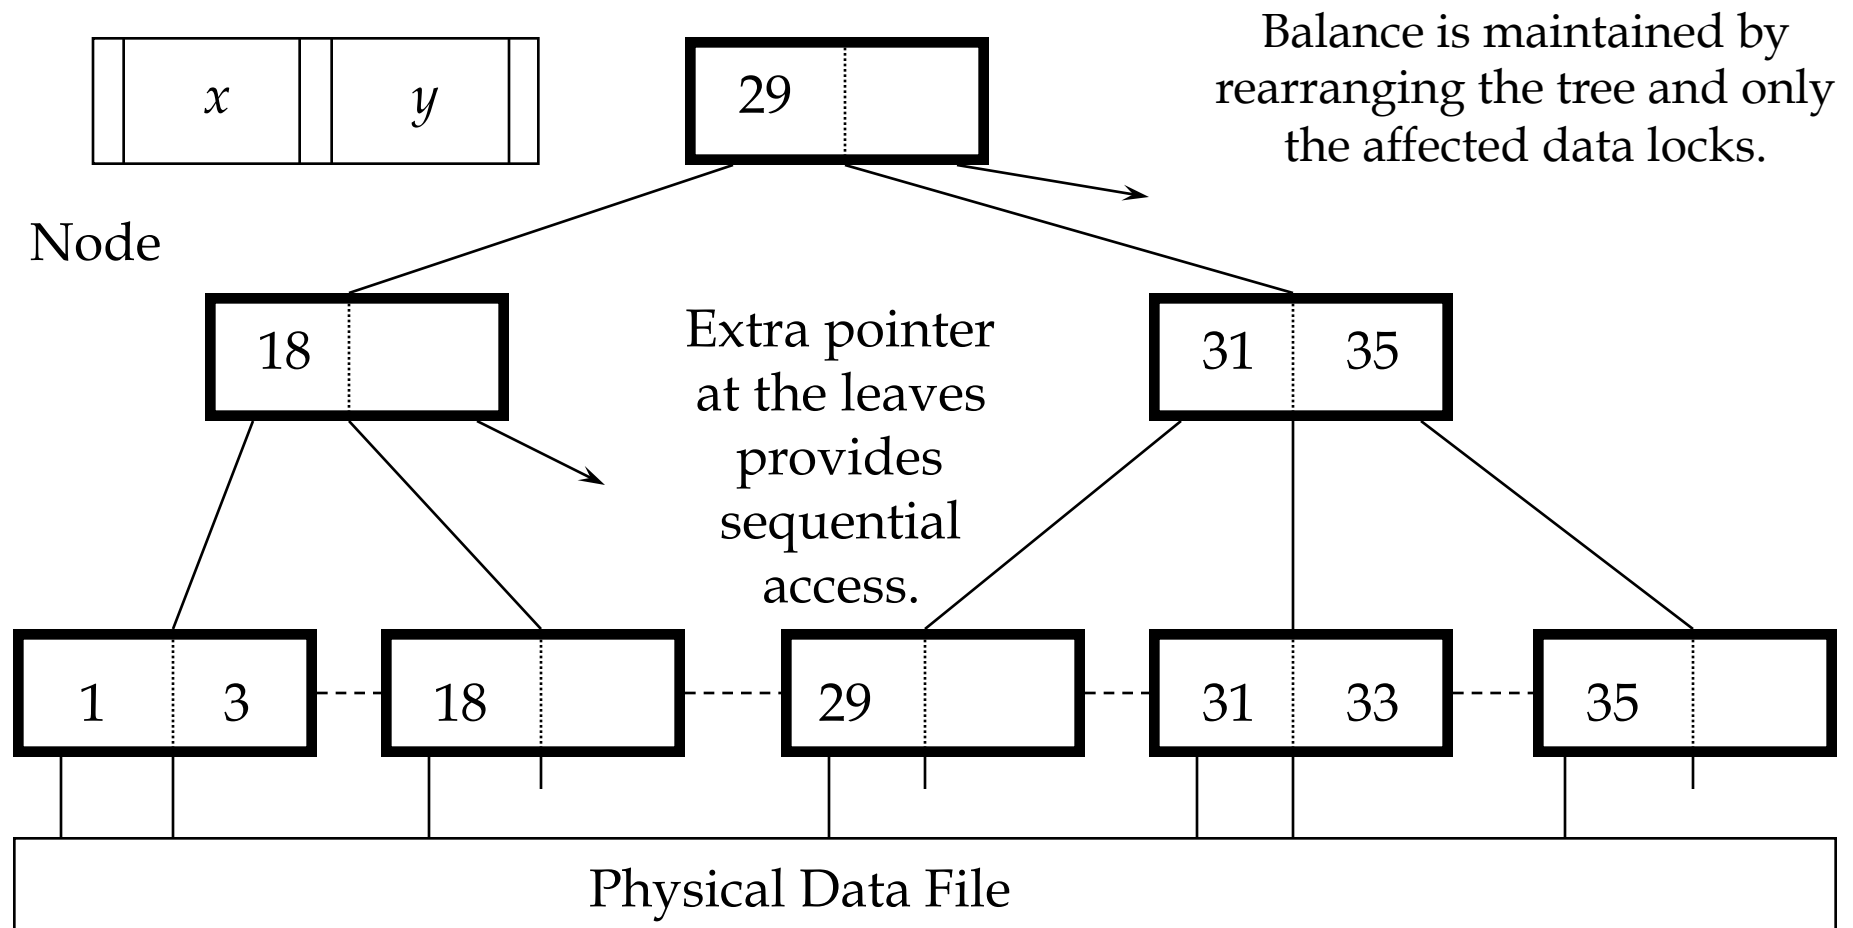
# Balanced B-Tree

- All paths from the root to a leaf node are the same length
  - Predictable access times
- Each node that is not a leaf has least  $n/2$  and at most  $n$  children, where  $n$  is the order of the B-tree.
- Leaf nodes contain at least  $(n-1)/2$  and at most  $n-1$  pointers to data record locations.

# B-Tree Index



# B+ Trees



# Reverse Key Index (B-Tree)

Index on EMP (EMPNO)

	KEY	ROWID
EMPNO	(BLOCK#	ROW# FILE#)
-----	-----	-----
1257	0000000F.0002.0001	
2877	0000000F.0006.0001	
4567	0000000F.0004.0001	
6657	0000000F.0003.0001	
8967	0000000F.0005.0001	
9637	0000000F.0001.0001	
9947	0000000F.0000.0001	
...	...	...
...	...	...

EMP table

EMPNO	ENAME	JOB	...
-----	-----	-----	...
7499	ALLEN	SALESMAN	
7369	SMITH	CLERK	
7521	WARD	SALESMAN	...
7566	JONES	MANAGER	
7654	MARTIN	SALESMAN	
7698	BLAKE	MANAGER	
7782	CLARK	MANAGER	
...	...	...	...
...	...	...	...

*What about range queries?*

# Use of B-Tree Indices

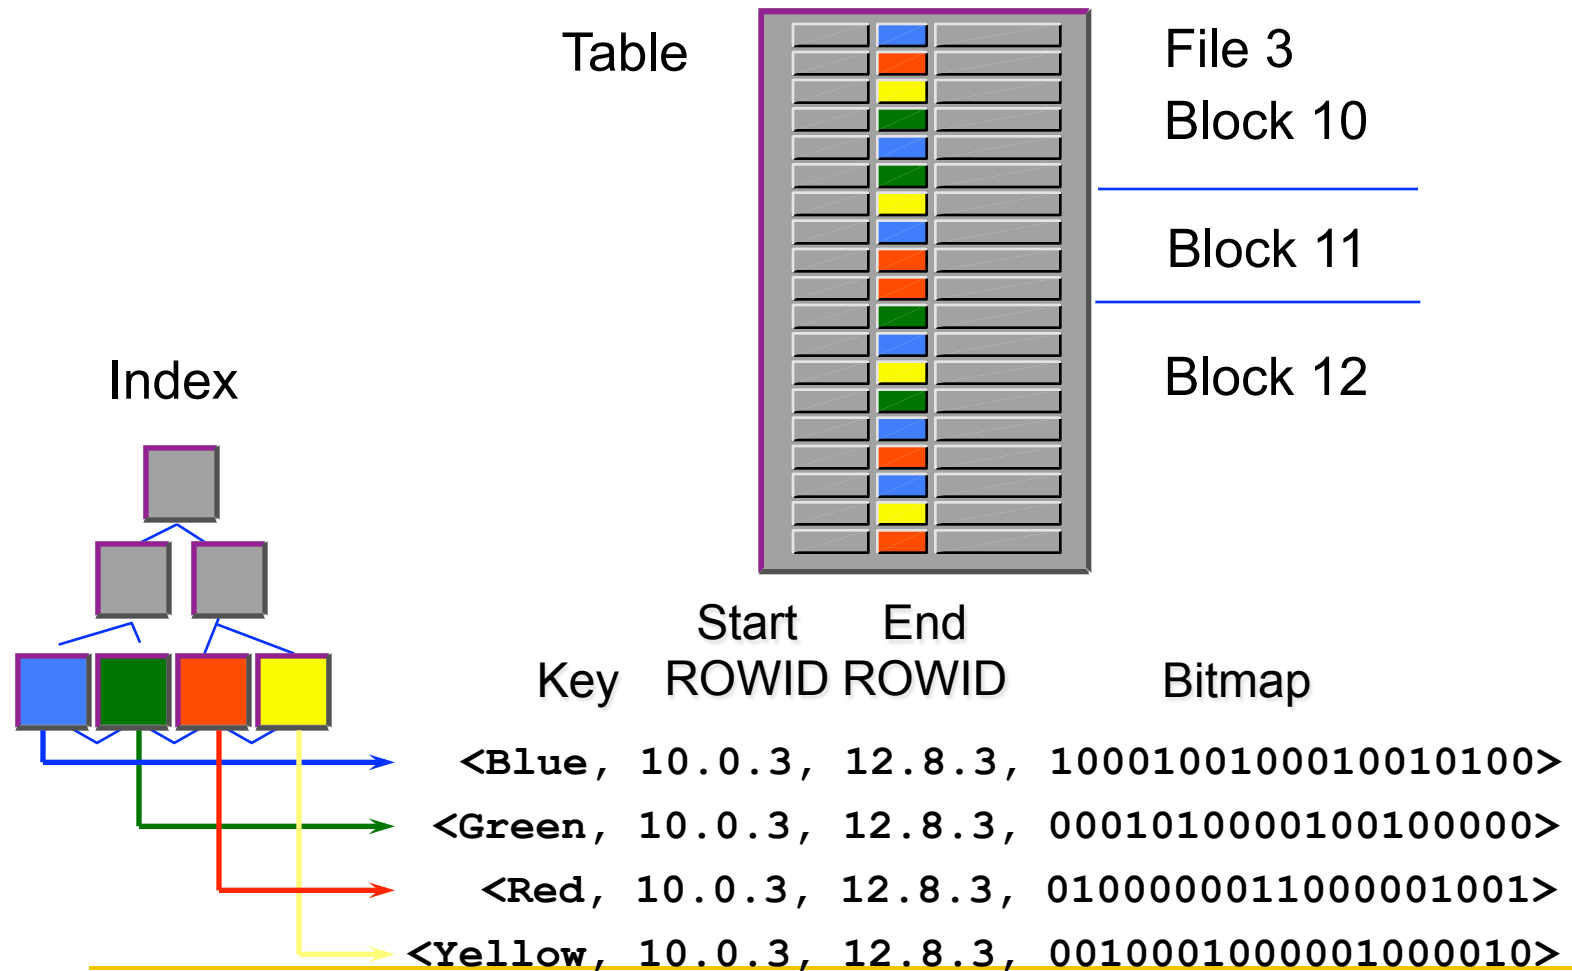
- Particularly suited to high-cardinality attributes
- The B-tree index is the default in most cases
- “For each primary key and unique constraint, Oracle automatically (implicitly) creates a B-tree index.”
- B-tree indexes use a large amount of space
- Index can be larger than data table

# Bitmap Indices

- A newer form of indexing that is very appropriate for data warehouses
- Appropriate for attributes with low cardinality (e.g., < 1%)
  - Examples: gender, Y/N attributes, categorical attributes ...



# Bitmap Index



# B-Tree vs Bitmap

B-Tree	Bitmap
Suitable for high-cardinality columns	Suitable for low-cardinality columns
Updates on keys relatively inexpensive	Updates to key columns very expensive
Inefficient for queries using OR predicates	Efficient for queries using OR predicates
Useful for OLTP	Useful for DSS

# Composite Indices

# Composite

- Given two attributes A1 and A2 ...
- If one has an index (idx\_a1) and the other does not, we can access the index first and then search the selected rows for a specific value of A2
- i.e. idx\_last\_name

```
SELECT last_name, first_name  
FROM Employee  
WHERE last_name = 'Smith' AND  
first_name LIKE 'F%';
```

# Composite

- If both A1 and A2 have individual indexes, we could use each index to retrieve the appropriate rows and then compute the intersection

idx\_last\_name and idx\_first\_name

- A composite index on A1 and A2 could be used to directly retrieve the required rows, but it is more query specific

idx\_last\_first\_name



# Administration

# Create Index

```
CREATE INDEX scott.emp_lname_idx  
ON scott.employees(last_name)  
PCTFREE 30  
STORAGE(INITIAL 200K NEXT 200K  
PCTINCREASE 0 MAXEXTENTS 50)  
TABLESPACE indx01;
```

Free space left in block  
before new is created

Keep indices in  
separate tablespace

# Rebuilding Indices

Use this command to:

- Move an index to a different tablespace
- Improve space utilization by removing deleted entries
- Change a reverse key index to a normal B-tree index and vice versa

```
ALTER INDEX scott.ord_region_id_idx REBUILD  
TABLESPACE indx02;
```



# Dropping Indices

- Drop and re-create an index before bulk loads
- Drop indexes that are infrequently needed and build them when necessary
- Drop and recreate invalid indexes

```
DROP INDEX scott.dept_dname_idx;
```

---

# Indexing Strategy

# Indexing Strategy

- Most RDBMS create an index on PKs
- Specify indices for *foreign keys* that are used for joining tables
- Index on *nonkey attributes* that are used frequently for selection criteria or grouping
- OLAP indexed more than OLTP
- Proactive – anticipating usage and build accordingly
- Reactive – based on optimizer and query implementation plan
- Too many on OLTP can slow updates
- Too few on OLAP can slow queries

# Oracle Indexing Guidelines

- Balance query and DML needs
- Place in separate tablespace
  - Old – Maximize I/O while minimize disk accesses
- Use uniform extent sizes: multiples of five blocks or MINIMUM EXTENT size for tablespace.
- Consider NOLOGGING for large indices.
- Set high PCTFREE if new key values are likely to be within the current range.

# Query Explain Plan

# Query Explain

- Each DBMS has a query optimizer
- Before a query is run, query optimizer develops an execution plan
  - Which columns used as index keys, have unique values
  - How many rows each table has
- Tools and statements help us view the execution plan and change queries accordingly
  - SQL Developer write the query and click:



# Query Explain Examples

**SELECT** employee\_id, last\_name, first\_name  
**FROM** employees  
**WHERE** last\_name = 'Feeney' **AND**  
 first\_name = 'Kevin';

## ■ No index on last\_name, first\_name

OPERATION	OBJECT_NAME	COST	LAST_CR_BUFFER_GETS
SELECT STATEMENT		3	
TABLE ACCESS (FULL)	EMPLOYEES	3	4
Filter Predicates			
AND			
LAST_NAME='Feeney'			
FIRST_NAME='Kevin'			

## ■ Index on last\_name, first\_name

OPERATION	OBJECT_NAME	COST	LAST_CR_BUFFER_GETS
SELECT STATEMENT		2	
TABLE ACCESS (BY INDEX ROWID)	EMPLOYEES	2	2
INDEX (RANGE SCAN)	EMP_NAME_IX	1	1
Access Predicates			
AND			
LAST_NAME='Feeney'			
FIRST_NAME='Kevin'			

# Query Explain Examples

```
SELECT employee_id, last_name
FROM employees JOIN jobs USING (job_id)
WHERE job_title = 'Stock Clerk'
ORDER BY last_name;
```

- No index on job\_title; index on job\_id (employees & jobs)
  - Total Cost = 5

OPERATION	OBJECT_NAME	COST	LAST_CR_BUFFER_GETS
SELECT STATEMENT		5	
SORT (ORDER BY)		5	5
NESTED LOOPS		5	5
NESTED LOOPS		4	4
TABLE ACCESS (FULL)	JOBS	3	3
Filter Predicates			
JOBS.JOB_TITLE='Stock Clerk'			
INDEX (RANGE SCAN)	EMP_JOB_IX	0	1
Access Predicates			
EMPLOYEES.JOB_ID=JOBS.JOB_			
TABLE ACCESS (BY INDEX ROWID)	EMPLOYEES	1	1



# Query Explain Examples Con't

- Index on job\_title and job\_id (employees & jobs)
  - Total Cost = 4

OPERATION	OBJECT_NAME	COST	LAST_CR_BUFFER_GETS
SELECT STATEMENT		4	
SORT (ORDER BY)		4	4
NESTED LOOPS		3	4
NESTED LOOPS		3	3
TABLE ACCESS (BY INDEX ROWID)	JOBS	2	2
INDEX (RANGE SCAN)	JOBS_JOB_TITLE_IX	1	1
Access Predicates			
JOBS.JOB_TITLE='Stock Cle			
INDEX (RANGE SCAN)	EMP_JOB_IX	0	1
Access Predicates			
EMPLOYEES.JOB_ID=JOBS.JOB_			
TABLE ACCESS (BY INDEX ROWID)	EMPLOYEES	1	1